

RAG-Based Clinical Decision Support / Triage Assistant

Complete Capstone Project Document

1. Literature Review (2020–2026)

Foundational RAG (2020). Lewis et al. (2020, NeurIPS) introduced retrieval-augmented generation for knowledge-intensive NLP tasks — the base architecture (retriever + generator) nearly all later clinical-RAG work extends.

Hallucination as the core motivating problem (2023–2024). Clinical-LLM literature is unusually explicit about the stakes here: <cite index="26-1">prior studies reported hallucination rates of 28.6–39.6% in medical-specialized models</cite> without retrieval grounding, and a JAMA Network Open study specifically examined accuracy/consistency/hallucination when LLMs analyze unstructured clinical notes. This is the direct justification for why RAG (not a bare LLM) is the right architecture for clinical decision support.

Domain-grounded RAG systems (2024). Multiple guideline-grounded systems emerged across specialties: a hepatology-guideline RAG framework (Krešević et al., *npj Digital Medicine*, 2024), and later ophthalmic and radiology-specific variants. These consistently show the same pattern — narrow, specialty-specific grounding reduces hallucination substantially but each system is built and evaluated in isolation with inconsistent methodology.

Architecture sophistication (2024–2025). Recent variants extend beyond plain dense retrieval: <cite index="24-1">hybrid RAG combining dense and sparse retrieval with reciprocal rank fusion, and knowledge-graph-enhanced retrieval over UMLS (a ~3M-node clinical concept graph) with graph-attention-network-based relation scoring</cite> represent the current state of the art in retrieval sophistication for this domain. Privacy-preserving, on-premises deployment variants (e.g., "MediRAG") specifically address the fact that clinical data often cannot leave institutional infrastructure — <cite index="24-1">though this vital privacy protection introduces additional system complexity.</cite>

Quantified hallucination-reduction results (2025). A radiology contrast-media consultation study found <cite index="26-1">complete elimination of hallucinations (reducing incidence from 8% to 0%) when using a RAG-enhanced local LLM</cite> compared to the ungrounded baseline — a strong, citable, concrete result for your own evaluation framing.

RAG vs. fine-tuning debate (2025). A systematic review (Jan 2023–Dec 2025 search window) concluded RAG is generally the more resource-efficient and scalable approach for medical LLM deployment, since fine-tuning requires substantial compute and annotated data often unavailable in clinical settings — though several studies found

hybrid RAG+fine-tuning approaches enhance accuracy further in specialized diagnostic tasks.

Evaluation methodology remains the weak point (2024-2026). This is your most important gap: work on "Ranking Over Scoring" (COLING 2025) argues for more reliable, robust automated evaluation of LLM-generated medical explanations, and a broader trustworthy-RAG survey (2025) catalogs LLM-as-judge as an emerging but still-inconsistent evaluation practice. Across nearly every paper reviewed, hallucination-rate measurement methodology differs (human review vs. automated fact-checking vs. LLM-as-judge), making cross-paper comparison unreliable.

Summary table

Period	Focus	Key contribution
2020	Foundational architecture	RAG introduced for knowledge-intensive NLP (Lewis et al.)
2023-2024	Problem motivation	Hallucination rates of 28.6-39.6% documented in ungrounded medical LLMs
2024	Domain-grounded systems	Specialty-specific guideline-grounded RAG (hepatology, ophthalmology)
2024-2025	Architecture sophistication	Hybrid dense+sparse retrieval; knowledge-graph-enhanced (UMLS) retrieval
2025	Quantified results	RAG reduces hallucination from 8% to 0% in a radiology-consultation study
2025	RAG vs fine-tuning	RAG favored for resource-constrained clinical deployment; hybrid approaches gaining traction
2024-2026	Evaluation gap	No standardized hallucination-measurement methodology across studies — your opening

2. Research Gap Analysis

- No standardized, reproducible hallucination-evaluation protocol.** Every paper measures hallucination differently (human review, automated fact-checking, LLM-as-judge), making results incomparable — this is explicitly flagged across multiple 2024-2026 sources above.
- Confidence-calibrated abstention is underexplored.** Most systems always answer; very few explicitly detect low-retrieval-confidence situations and abstain or flag uncertainty — a real patient-safety gap, since a confidently wrong answer is more dangerous than a system that says "I'm not sure, consult a specialist."

3. **Narrow, single-specialty focus limits generalizable conclusions.** Existing systems (hepatology, ophthalmology, radiology) are each built and evaluated in isolation — a cross-domain evaluation using one consistent methodology hasn't been done at the capstone-accessible level.

Your novelty claim should combine gaps #1 and #2: a guideline-grounded RAG assistant with confidence-calibrated abstention, evaluated using a rigorous, reproducible hallucination-measurement protocol (combining automated fact-checking against retrieved source passages + structured human review) — directly addressing the methodology-inconsistency gap while adding a genuine safety mechanism (abstention) that most reviewed systems lack.

3. Existing Solutions Comparison

System/Approach	Strength	Weakness (your opening)
Vanilla LLM (no retrieval)	Simple, fast	28.6-39.6% hallucination rate in medical-specialized models — unsafe for clinical use
Hepatology/Ophthalmology guideline-RAG (2024)	Strong specialty-specific hallucination reduction	Narrow scope, isolated evaluation methodology, no abstention mechanism
Hybrid dense+sparse + UMLS knowledge-graph RAG	State-of-the-art retrieval sophistication	Complex to build in 3-4 months; doesn't address the evaluation-standardization or abstention gaps
MediRAG (privacy-preserving, on-premises)	Addresses real data-residency constraint	Added complexity; privacy focus, not evaluation-rigor or abstention focus
Fine-tuned medical LLM (no RAG)	Can outperform RAG on narrow diagnostic tasks in hybrid setups	Requires substantial compute/annotated data — impractical for a 3-4 month capstone

Recommendation: Build a **dense-retrieval RAG pipeline** (skip the full knowledge-graph complexity — it's not your novelty axis and would eat your timeline) using a vector DB (e.g., pgvector or Chroma) over a curated guideline corpus, and invest your saved time into the **abstention mechanism and evaluation protocol**, which is where your actual contribution lives.

4. Novel Contribution (your paper's thesis)

"Confidence-Calibrated Retrieval-Augmented Generation for Safe Clinical Decision Support: A Reproducible Hallucination-Evaluation Protocol"


A guideline-grounded RAG system that (a) explicitly scores retrieval confidence per query and abstains/flags uncertainty below a threshold rather than always answering, and (b) is evaluated using a documented, reproducible protocol combining automated citation-grounding checks (does the answer's claims trace back to retrieved passages) with structured human review — directly addressing the standardization gap the literature repeatedly flags, while adding a safety mechanism (abstention) most existing systems don't have.


5. Patentability Analysis

Low patentability — this should be a pure publication play, and that's fine.


- RAG architecture itself, and even confidence-thresholded abstention, are not novel *mechanisms* in the abstract (confidence thresholding exists broadly in ML) — the novelty here is in application and evaluation rigor, which is publishable but not independently patent-eligible.
 - **No realistic patent path here** — don't spend capstone time on IP filing for this one. Focus entirely on making the evaluation methodology rigorous enough to be genuinely citable; that's your real asset.
-

6. Complete System Architecture





7. Database Schema (ER Diagram)



8. API Design

Method	Endpoint	Purpose
POST	<code>/auth/login</code>	User authentication (JWT)
POST	<code>/guidelines/upload</code>	Ingest a new guideline document into the corpus
POST	<code>/query</code>	Submit a clinical question, get grounded answer + citations or abstention
GET	<code>/query/{id}</code>	Retrieve a past query's full response + retrieved chunks
POST	<code>/evaluation/run</code>	Trigger an evaluation run against a benchmark (e.g., MedQA subset)
GET	<code>/evaluation/{id}/results</code>	Fetch hallucination-rate, abstention-rate, citation-accuracy results
POST	<code>/review/{response_id}</code>	Submit human-review annotation (hallucination flag, notes)
GET	<code>/corpus/stats</code>	Corpus size, specialty coverage, chunk count

9. Frontend Wireframes

Clinical RAG Assistant

[Login: clinician]

Ask a question:

"What's the recommended first-line treatment for..."

[Submit]

Answer (confidence: 0.87 – high)

"Based on [Guideline X, Section 3.2]..."

Sources: [1] Guideline X [2] Guideline Y

--- OR, low-confidence case ---

▲ Confidence: 0.31 – Insufficient grounding found.


This question may fall outside the current guideline corpus. Please consult a specialist.

[Researcher view] Evaluation Dashboard

Hallucination rate: 4.2% Abstention rate: 11%

Citation accuracy: 93%

10. Backend Microservice Architecture



Modular monolith is again the right call here for 2 builders — separate Query, Retrieval, Confidence, and Evaluation as internal modules/services within one deployable app rather than true separate microservices.

11. AI/ML Model Design

- **Retriever:** dense embedding retrieval (e.g., a strong open embedding model) over chunked guideline text, stored in pgvector or Chroma.
- **Generator:** an LLM API call with retrieved chunks injected into the prompt, explicitly instructed to cite which chunk(s) support each claim (this citation requirement is what enables your citation-grounding check).
- **Confidence scorer:** combine (a) top-k retrieval similarity scores and (b) an LLM-as-judge or entailment-model check of whether the generated answer's claims are actually supported by the retrieved chunks — abstain if either signal is below threshold.
- **Citation-grounding checker:** for each claim in the answer, verify it can be traced to a specific retrieved chunk (simple entailment check or exact/paraphrase match against chunk content) — this is your automated half of the hallucination-evaluation protocol.
- **Threshold tuning:** treat the abstention threshold as an experimental variable — report the precision/coverage tradeoff (higher threshold → fewer hallucinations but more abstentions) as a curve, not a single fixed number.

12. Dataset Recommendations

Dataset/Corpus	Use	Why it fits
Curated public clinical guidelines (e.g., WHO, NICE, or specialty-society guidelines, whichever are freely redistributable)	Your grounding corpus	Directly matches the "guideline-grounded" framing used in the 2024 literature you're extending
MedQA / PubMedQA	Evaluation benchmark	Standard, widely used medical QA benchmarks — enables direct comparison to prior RAG-vs-vanilla-LLM results
MedHALT (Med-Halt hallucination test, referenced above)	Hallucination-specific evaluation	Purpose-built for measuring medical hallucination — strengthens your evaluation-rigor claim
A held-out "out-of-corpus" question set (questions outside your guideline corpus's scope)	Abstention evaluation	Essential for measuring whether your system correctly abstains rather than hallucinating on unfamiliar topics

Scope note: pick ONE clinical sub-domain (e.g., primary-care triage, or a single specialty) rather than trying to cover all of medicine — this keeps your guideline corpus

curatable in the available time and matches how the existing literature is (necessarily) specialty-scoped.

13. Evaluation Metrics

- **Hallucination rate:** % of answered (non-abstained) responses containing at least one unsupported claim — measured via your automated citation-grounding checker, spot-validated by human review
 - **Abstention rate:** % of queries where the system declines to answer — report alongside hallucination rate as a tradeoff curve, not in isolation
 - **Citation accuracy:** % of cited sources that genuinely support the claim they're attached to
 - **Coverage-accuracy tradeoff:** as you raise the confidence threshold, hallucination should drop but abstention should rise — plot this curve, it's your key result
 - **Comparison baselines:** vanilla LLM (no retrieval), RAG without abstention (always answers), your full system — report hallucination rate for all three
 - **Human-agreement check:** inter-rater agreement between your automated citation-checker and human reviewers on a sample — validates your automated metric is trustworthy, addressing the evaluation-standardization gap directly
-

14. Week-by-Week Implementation Plan (16 weeks)

Week	Milestone
1	Literature review finalized; pick clinical sub-domain; team roles locked
2	Guideline corpus curated and chunked; vector DB set up
3	Basic dense retrieval pipeline working (query → top-k chunks)
4	LLM generation with citation-instructed prompting integrated
5	Baseline vanilla-LLM (no retrieval) comparison system built for later contrast
6	Confidence scoring module v1 (retrieval-similarity based)
7	Citation-grounding checker v1 (automated claim-to-chunk verification)
8	Abstention mechanism wired end-to-end; internal review 1 prep
9	Internal Review 1 — present pipeline + preliminary hallucination-rate numbers
10	Confidence scorer v2: add LLM-as-judge or entailment-model signal, not just retrieval similarity
11	Evaluation harness built against MedQA/PubMedQA/MedHALT subsets
12	Full evaluation run: vanilla LLM vs RAG-no-abstain vs full system; threshold sweep for tradeoff curve
13	Human-review sample for inter-rater validation of automated metric; dashboard built
14	Internal Review 2 — present full results; incorporate feedback
15	Paper: Results/Discussion finalized; target-venue formatting
16	External Review 3; submission to target venue

15. Team Task Allocation (2 developers + 1 researcher)

Developer A — Retrieval & Corpus

- Weeks 1-4: guideline corpus curation, chunking strategy, vector DB setup, retrieval pipeline
- Weeks 5-10: confidence scoring (retrieval-similarity component), corpus-management API
- Weeks 11-16: dashboard backend, evaluation-harness infrastructure, deployment/demo

Developer B — Generation & Safety Mechanism

- Weeks 1-4: LLM generation integration, citation-instructed prompting
- Weeks 5-10: citation-grounding checker (this is core to your novelty — prioritize it), abstention logic, LLM-as-judge confidence signal
- Weeks 11-16: full evaluation-run execution, threshold-sweep analysis, results processing

Researcher — Literature, Writing, Evaluation Design

- Weeks 1-2: literature review, gap analysis, finalize novelty framing (evaluation-standardization + abstention)
- Weeks 3-10: draft Methods in parallel with build; design the human-review protocol (inter-rater agreement setup) *with* the developers so it's actually executable at week 13
- Weeks 11-16: Results, Discussion, Abstract, Introduction; venue formatting; review-panel prep; personally run/coordinate the human-review sample since this needs careful, unbiased protocol design

16. IEEE / ACM / Springer Conference Recommendations

- **IEEE conferences** on Healthcare Informatics, Biomedical & Health Informatics (e.g., IEEE regional health-AI symposiums) — standard fit.
- **ACM CHIL-adjacent workshops** or ACM regional health-informatics symposiums — directly relevant to clinical NLP/RAG work.
- **Springer LNCS/LNAI volumes** on medical NLP or clinical AI — accessible for first-time student authors.
- **JMIR AI** and similar venues have published exactly this kind of clinical-RAG evaluation work recently (see literature review above) — worth checking their current submission scope even if primarily a journal rather than conference track, since a strong capstone result could be journal-appropriate too.

Action item: verify live CFPs/deadlines 2-3 months before your week-15 target.

17. Resume Impact Analysis

Resume bullet (example):

"Built a guideline-grounded RAG clinical assistant with confidence-calibrated abstention, reducing hallucination rate to X% (vs. Y% for ungrounded baseline) using a reproducible evaluation protocol combining automated citation-grounding and human review."

Why this lands with recruiters in 2026:

- RAG/LLM-grounding skills are currently the single hottest applied-AI competency in job postings — this project speaks directly to that demand.
- "Confidence-calibrated abstention" signals safety-conscious AI engineering, a differentiator as companies increasingly care about responsible-AI practices, not

just raw capability.

- A rigorous, documented evaluation methodology (not just "I built a chatbot") is exactly what separates a genuinely interview-worthy AI project from a weekend hackathon demo.
 - Fast to build given your existing RAG experience, which means more of your 3–4 months goes into the evaluation rigor that actually makes this resume-differentiating.
-

18. Interview Questions Based on This Project

1. Why RAG instead of fine-tuning for this use case — what's the actual tradeoff?
 2. How does your confidence scorer decide when to abstain, and how did you pick the threshold?
 3. Walk through your citation-grounding checker — how do you verify a claim is actually supported by a retrieved chunk, and where does it fail?
 4. What's the difference between retrieval-similarity confidence and an LLM-as-judge confidence signal — why use both?
 5. How did you validate that your automated hallucination metric is trustworthy (inter-rater agreement with human review)?
 6. What happens when the guideline corpus itself is outdated or wrong — how would you handle that failure mode?
 7. How would you extend this system to handle multi-turn clinical conversations rather than single-shot Q&A?
 8. What are the real deployment blockers between this prototype and actual clinical use (regulatory approval, liability, integration with EHR systems)?
-

19. Future Startup / Commercialization Opportunities

- **B2B health-AI infrastructure play:** license the confidence-calibrated RAG + evaluation framework to health-AI vendors or hospital systems building their own clinical assistants — the evaluation rigor itself (not just the chatbot) is a sellable asset, since the field's biggest documented weakness is inconsistent evaluation.
 - **Realistic first customer:** medical education platforms or clinical-decision-support vendors who need a safer, better-evaluated alternative to a bare LLM wrapper — lower regulatory burden than direct patient-facing deployment.
 - **Differentiator:** most current commercial "AI medical assistants" don't publicly report hallucination rates or have abstention mechanisms — genuinely leading with measured safety is a real trust/marketing wedge, not just a technical one.
 - **Caution:** direct clinical deployment requires regulatory clearance (e.g., as a Software as a Medical Device in various jurisdictions) — realistic path is "medical-education or clinician-support tool with a multi-year path to regulated deployment," not an immediate consumer product.
-

All three projects (FL+DP, RL Microgrid, RAG Clinical) are now fully scoped at this depth. Want a final side-by-side comparison to help you commit to one, or should we start turning your top choice's week-1 tasks into an actual checklist?